

## TP 1 : Estimators – Stochastic gradient descent

### Exercise 1 : Comparing different estimators for the uniform model (10pt)

We consider the parametric family of uniform laws  $\{\mathcal{U}(\theta), \theta \in \mathbb{R}\}$  on  $\mathbb{R}$  where  $\mathcal{U}(\theta)$  admits the following density with regard to the Lebesgue's measure :

$$p_\theta(x) = \frac{1}{\theta} \mathbb{1}_{[0,\theta]}.$$

We consider independent random samples  $X_1, X_2, \dots, X_n$  of density  $p_\theta$ . The goal of this exercise is to compare different estimators for the parameter  $\theta$ .

1. Calculate  $\mathbb{E}_\theta(X_1)$  and deduce an estimator  $\hat{\theta}_1$  of  $\theta$  using the method of moments. (1pt)
2. Calculate the quadratic risk of  $\hat{\theta}_1$ . (3pt)
3. Calculate the maximum likelihood estimator  $\hat{\theta}_2$ . (2pt)
4. Calculate the quadratic risk of  $\hat{\theta}_2$ . (3pt)
5. Which estimator is preferable ? (1 pt)

### Exercise 2 : Box-Muller and Marsaglia-Bray algorithm (Bonus +3pt)

Let  $R$  a random variable with Rayleigh distribution with parameter 1 and  $\Theta$  with uniform distribution on  $[0, 2\pi]$ . We also assume that  $R$  and  $\Theta$  are independent. The density of  $R$  with regard to the Lebesgue's measure writes :

$$\forall r \in \mathbb{R}, \quad f_R(r) = r \exp\left(-\frac{r^2}{2}\right) \mathbb{1}_{\mathbb{R}^+}(r)$$

1. Let  $X$  and  $Y$  such that

$$X = R \cos(\Theta) \quad \text{and} \quad Y = R \sin(\Theta).$$

Prove that both  $X$  and  $Y$  have  $\mathcal{N}(0, 1)$  distribution and are independent.

2. Write an algorithm for sampling independent Gaussian distribution  $\mathcal{N}(0, 1)$ .
3. Consider the Marsaglia-Bray algorithm given below.
  - a) What is the distribution of  $(V_1, V_2)$  at the end of the "while" loop ?
  - b) What is the expected number of steps in the "while" loop ?

**Algorithm 1:** Marsaglia-Bray algorithm

---

```

1  $V_1 = 1, V_2 = 1$  while  $V_1^2 + V_2^2 > 1$  do
2   | Sample  $U_1, U_2$  independant r.v. with distribution  $\mathcal{U}([0, 1])$  ;
3   | Set  $V_1 = 2U_1 - 1$  and  $V_2 = 2U_2 - 1$ .
4 end
5 Set  $S = \sqrt{-2 \log(V_1^2 + V_2^2)}$  ;
6 Set  $X = S \frac{V_1}{\sqrt{V_1^2 + V_2^2}}$  and  $Y = S \frac{V_2}{\sqrt{V_1^2 + V_2^2}}$  ;
7 return  $(X, Y)$ .

```

---

c) Set

$$T_1 = \frac{V_1}{\sqrt{V_1^2 + V_2^2}}, \quad \text{and} \quad V = V_1^2 + V_2^2.$$

Show that  $T_1$  and  $V$  are independent,  $V \sim \mathcal{U}([0, 1])$  and  $T_1$  has the same distribution as  $\cos(\Theta)$  with  $\Theta \sim \mathcal{U}([0, 2\pi])$ .

d) What is the distribution of the output  $(X, Y)$  ?

**Exercise 3 : Learning a linear classifier, [Bot91, BCN16] (10 pt)**

We consider the supervised learning problem of assigning inputs  $x \in \mathbb{R}^d$  to one of two categories labeled by  $y \in \{-1, +1\}$ . The statistical relationship between inputs and labels is governed by an unknown joint distribution  $\mathbb{P}(x, y)$ . If this distribution were known, classification could be optimally performed by computing the conditional probability  $\mathbb{P}(y | x)$  and applying the Bayes decision rule.

In practice, learning refers to the process of inferring this decision mechanism from a finite collection of observations. This dataset is composed of  $n$  labeled samples :

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\},$$

where each  $x_i \in \mathbb{R}^d$  denotes a feature vector, and  $y_i \in \{-1, +1\}$  indicates its corresponding class label. The goal is to construct a classification model  $h(x)$  that performs well not just on these training examples but also on unseen data — that is, it should generalize.

In this exercise, we only consider **linear classifiers**, that is functions of the form

$$h(x; w, \tau) = \langle w, x \rangle + \tau,$$

where  $w \in \mathbb{R}^d$  is a weight vector and  $\tau \in \mathbb{R}$  is a bias term. This model attempts to separate the two classes using a hyperplane in the feature space. The bias  $\tau$  can be adjusted to favor either precision ( $\mathbb{P}[y = 1 | h(x) = 1]$ ) or recall ( $\mathbb{P}[h(x) = 1 | y = 1]$ ) depending on application needs.

Prediction is done via the sign of  $h(x)$ , with the predicted label being  $\text{sign}(h(x))$ . However, optimizing such a model directly using classification accuracy is computationally challenging due to the non-differentiability of the sign function. To overcome this, a **continuous loss function** is introduced to approximate the classification objective and facilitate optimization.

### The Adaline Model

One early and influential model is the **Adaline** (Adaptive Linear Neuron), proposed by Widrow and Hoff in 1960. It employs a linear prediction function :

$$h(x; w, \tau) = \langle w, x \rangle + \tau,$$

For convenience, we can consider the extended feature vector  $\bar{x} = (x, 1)$  in order to regroup  $\tau, w$  in one parameter  $\bar{w}$ . Then  $h(\bar{x}, \bar{w}) = \langle \bar{w}, \bar{x} \rangle$ . The Adaline model evaluates the classifier's performance using the **mean squared error (MSE)** :

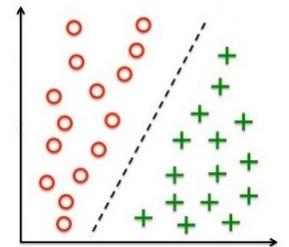
$$L(w) = \mathbb{E}_{(x,y) \sim \mathbb{P}} [(y - h(\bar{x}, \bar{w}))^2].$$

However, since  $\mathbb{P}$  is unknown, we cannot compute this quantity exactly. Instead, we approximate it using the available dataset. This leads to the **empirical risk**, defined as :

$$L_n(w) = \frac{1}{n} \sum_{i=1}^n (y_i - h(\bar{x}_i, \bar{w}))^2.$$

Minimizing  $L_n(w)$  provides an estimate of the true minimizer of  $L(w)$ , under the assumption that the training data is representative of the underlying distribution.

- (Python) Generate a set of observations  $(\{(x_i, y_i)\}_{i=1}^N$  by sampling random points  $x_i$  in  $\mathbb{R}^2$  and choosing an arbitrary normal vector  $w^*$  to define the separating hyperplane of the two classes. The sign of  $\langle w^*, x_i \rangle$  gives the label  $y_i \in \{-1, +1\}$ . (2pt)



- Write an algorithm for minimizing  $L_n(w)$  with Stochastic Gradient Descent.(2pt)
- (Python) Implement and test the algorithm on the set of observations generated at the first question. What is the vector  $\hat{w}$  estimated ? Is it far from  $w^*$  ? (2pt)

4. (Python) Noise your observations  $\{z_i\}_{i=1}^n$  with an additive Gaussian noise and perform the optimisation again. Compare with the result of question three.(1pt)
5. (Python) Test the algorithm on the *Heart Disease (Diagnostic) Data Set* [WSM95] : <https://archive.ics.uci.edu/dataset/45/heart+disease> . (3pt)

## Références

- [BCN16] Léon Bottou, Frank E. Curtis, and Jorge Nocedal. Optimization methods for large-scale machine learning. *eprint arXiv :1606.04838*, 2016.
- [Bot91] Léon Bottou. Stochastic gradient learning in neural networks. In *Neuro-Nîmes 91*, 1991.
- [WSM95] William H. Wolberg, W. Nick Street, and Olvi L. Mangasarian. UCI machine learning repository, 1995.